

Haris Faisal

haris.faisal@mail.utoronto.ca | (437)-665-1084 | [LinkedIn](#) | [GitHub](#)

EDUCATION

University of Toronto

Honors in Bachelor of Science, Computer Science Specialist + PEY Co-op,
Applied Statistics Major, Mathematics Minor Student

Sep 2022 – Dec 2026
Mississauga, Canada

- **Related Coursework:** Data Structures and Algorithms, Databases, Web Development, Computer Networks, Software Design, Machine Learning, Neural Networks, Operating Systems, Theory of Computation.

SKILLS

Programming: JavaScript, TypeScript, SQL, Python, HTML/CSS, Java, PostgreSQL, MongoDB, C, C++.

Tools & Frameworks: React, GraphQL, Node.js, Postman, Git, Django, FastAPI, AWS (SQS, Lambda, S3, EC2, ECS), Docker, Kubernetes, Vercel, Microsoft Office.

EXPERIENCE

BorderPass

Software Developer Intern

Sep 2025 – Present
Toronto, Canada

- Developed end-to-end onboarding and application workflows at a fast-growing legal tech startup using React, TypeScript, Node.js, and GraphQL, building reusable, extensible components for complex business workflows.
- Built an AI-powered eligibility assessment system using GenAI and BullMQ, offloading high-compute processing to background job queues to automate evaluations across 800+ applications.
- Automated manual eligibility checks previously handled by legal teams using the GenAI-powered assessment pipeline, reducing lawyer workload by ~50% and accelerating application turnaround time.
- Designed and integrated AI-driven document parsing pipelines using AWS Lambda and S3, building a scalable integration architecture for structured extraction of application documents at scale.
- Shipped reliable, resilient production features through a CI/CD pipeline with unit and integration testing, collaborating with product, design, and legal teams to optimize performance while deploying on AWS (EC2, ECS) and Vercel.

Denovonet

Software Engineering Intern

May 2025 – Aug 2025
Lahore, Pakistan

- Built and unit-tested Django backend APIs for rule-based anomaly detection, writing and optimizing SQL queries against large-scale Apache Iceberg tables via Trino for high-throughput, low-latency data access.
- Containerized asynchronous data processing and model-serving components using Docker and deployed on Kubernetes, enabling scalable, reproducible infrastructure for high-volume workflow automation.
- Collaborated cross-functionally with ML engineers to preprocess large transactional datasets for anomaly detection, improving alert precision and reducing false positives across distributed data pipelines.
- Implemented high-performance frontend dashboards in React using Vis.js and Recharts to visualize real-time alerts and transaction patterns, enhancing analyst workflows and enabling faster data-driven decisions.

PROJECTS

Mergent

Personal Project

July 2026 – Present
Toronto, Canada

- Building a FastAPI, PostgreSQL, and React application integrated with the Anthropic Claude API, to automate pull request review and generation for developers.
- Using Celery and Redis to run PR analysis and code generation as background tasks, keeping long-running AI processing separate from user-facing requests, so the system can handle pull requests without blocking users.
- Designing a multi-agent pipeline that explores codebases and generates context-aware code suggestions, coordinating hand-offs between agents, to reduce manual review overhead on repetitive PR tasks.

AI Powered Interactive Learning Platform

Computer Science Implementation Course, University of Toronto

May 2025 – Aug 2025
Remote

- Developed a full-stack AI-powered learning platform using Agile practices, with a React frontend featuring interactive visualizations for processor pipeline data, to support a personalized learning workflow for 300+ students.
- Implemented a GraphRAG system with Pinecone as the vector database and prompt engineering to retrieve relevant course context, delivering accurate, context-aware AI responses at scale.
- Designed and containerized RESTful API services using FastAPI and Docker to handle user state tracking, and chatbot synchronization with real-time on-screen content, ensuring reliable performance across environments.